Figure R2. Iterative SRD self-correction across diverse rare concepts. Three examples demonstrating the SRD convergence process. *(Top, DALL-E 3)* A complex mythological entity (Hindu deity Yama) requiring 3 rounds: initial RAVEL output exhibits missing arms, incorrect skin color, and wrong weapon; SRD R1 corrects skin tone but introduces extra heads; SRD R2 resolves arm count and buffalo vehicle; SRD R3 achieves full convergence (GSI>0.85) with correct crown, weapon, and posture. *(Bottom left, GLM-Image)* Hanfu converges in 1 round: initial output captures the rabbit and moon but lacks floral details; SRD R1 enlarges the moon and adds floral ornamentation, reaching threshold. *(Bottom right, GLM-Image)* Bleeding Tooth Fungus converges in 2 rounds: initial output captures shape and color but misses blood-drop texture and body pattern; SRD R1 fixes color and blood-drop morphology; SRD R2 achieves correct habitat, texture, and full structural accuracy. The attribute-level annotations illustrate how the stability matrix identifies and targets specific missing features at each iteration, enabling precise, non-destructive convergence rather than global image regeneration. (*Reference images are provided for comparison purposes only and are not used during generation.)